





Les 7: Semantic Image segmentation




Computer vision applications

Image Enhancement



image


Classification



“dog”

label

Embedding / Metric Learning



“45A4CE12006”

*description
vector*




Computer vision applications

Object detection



labeled
bounding
boxes

Semantic Segmentation



labeled pixels

Semantic Image Segmentation

- Semantic segmentation:
 - Classic techniques using colour
 - Fully Convolutional Network
 - Deconvolution
 - Skip connections: U-Net
- Instance Segmentation
 - Two-stage: Mask-RCNN
 - One-stage: YOLACT

SKIN COLOR SEGMENTATION



INPUT IMAGE



BINARY MASK

SKIN COLOR SEGMENTATION



BINARY MASK



BINARY MASK AFTER
PROCESSING


Semantic Image Segmentation

- Semantic segmentation:
 - Classic techniques using colour
 - Fully Convolutional Network
 - Deconvolution
 - Skip connections: U-Net
- Instance Segmentation
 - Two-stage: Mask-RCNN
 - One-stage: YOLACT


Semantic and instance segmentation



Image classification




Object detection



Semantic segmentation

"stuff", each **pixel** is classified




Instance segmentation

"things", each **instance** is classified

Semantic segmentation

- Each **pixel** is classified
- Output of network is tensor of shape **L x W x H**
 - With L the number of classes, and (W, H) the size of the input image
- E.g. U-Net
 - Ronneberger, O., Fischer, P., & Brox, T. (2015). U-Net: Convolutional Networks for Biomedical Image Segmentation. *ArXiv:1505.04597 [Cs]*. <http://arxiv.org/abs/1505.04597>




Instance segmentation

- Each **instance** is segmented and classified
 - Like object detection, but with segmentation masks instead of bounding boxes
- Output of the network is **a list of masks** with localization and classification predictions
- E.g. Mask R-CNN
 - He, K., Gkioxari, G., Dollár, P., & Girshick, R. (2017). Mask R-CNN. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2961–2969. <https://doi.org/10.1109/TPAMI.2018.2844175>



Extra: panoptic segmentation

(won't go into detail)



Semantic segmentation



Instance segmentation




Panoptic segmentation

"stuff" and "things": **each pixel**
and **each instance** is classified

Evaluation metric

- Pixel classification!
- Accuracy?
 - Heavily unbalanced
 - Common classes are over-emphasized
- *Intersection over Union*
 - Average across classes and images
- Per-class accuracy
 - Compute accuracy for every class and then average



Things vs Stuff

THINGS

- Person, cat, horse, etc
- Constrained shape
- Individual instances with separate identity
- May need to look at objects



STUFF

- Road, grass, sky etc
- Amorphous, no shape
- No notion of instances
- Can be done at pixel level
- “texture”




Challenges in data collection

- Precise localization is hard to annotate
- Annotating every pixel leads to heavy tails
- Common solution: annotate few classes (often things), mark rest as “Other”
- Common datasets: PASCAL VOC 2012 (~1500 images, 20 categories), COCO (~100k images, 20 categories)

Segmentation datasets

- Pascal Context
 - 10k+ training images with 400+ labels
 - semantic segmentation dataset



Segmentation datasets

- Microsoft Common Objects in COntext (COCO)
 - Originally: instance segmentation, but nowadays also contains annotations for stuff segmentation
 - 328k images with 91 labels



Semantic Image Segmentation


- Semantic segmentation:
 - Classic techniques using colour
 - Fully Convolutional Network
 - Deconvolution
 - Skip connections: U-Net
- Instance Segmentation
 - Two-stage: Mask-RCNN
 - One-stage: YOLACT

Fully Convolutional Network (FCN)


Long, J., Shelhamer, E., & Darrell, T. (2014). *Fully Convolutional Networks for Semantic Segmentation*. <https://arxiv.org/abs/1411.4038v2>

- Network for **semantic segmentation**
- Fully-connected layers in typical CNN classifiers are **replaced by conv layers**
 - Kernel size = size of fully-connected layer input
- Hence, network can take any input size
- Output feature maps are up-sampled to original input image size with **transposed convolutions**


A typical CNN classifier




Replace fully-connected layers with convolutions




Network can take any input size!




Upsample output feature maps




Semantic segmentation using convolutional networks




Semantic segmentation using convolutional networks




Semantic segmentation using convolutional networks



Semantic segmentation using convolutional networks




Semantic segmentation using convolutional networks



Semantic segmentation using convolutional networks

- Pass image through convolution and subsampling layers
- Final convolution with #classes outputs
- Get scores for *subsampled* image
- Upsample back to original size

Fully Convolutional Network (FCN)



Fully Convolutional Network





Semantic Image Segmentation

- Semantic segmentation:
 - Classic techniques using colour
 - Fully Convolutional Network
 - Deconvolution
 - Skip connections: U-Net
- Instance Segmentation
 - Two-stage: Mask-RCNN
 - One-stage: YOLACT



The resolution issue

- Problem: Need fine details!
- Shallower network / earlier layers?
 - Deeper networks work better: more abstract concepts
 - Shallower network => Not very semantic!
- Remove subsampling?
 - Subsampling allows later layers to capture larger and larger patterns
 - Without subsampling => Looks at only a small window!
- After subsampling: reverse operation = upsampling needed!

Reversing layers: Unpooling



Reversing layers: Deconvolution



Deconvolution = transposed convolution
(=up-convolution)

- <https://naokishibuya.medium.com/up-sampling-with-transposed-convolution-9ae4f2df52d0>

Transposed convolutions

- Aka up-convolution, fractionally-strided convolution or deconvolution
- Implemented by swapping forward and backward of regular convolution
- Transposed conv with stride s , kernel size k and padding p will produce output size of


$$o' = s(i' - 1) + k - 2p$$

Dumoulin, V., & Visin, F. (2018). A guide to convolution arithmetic for deep learning. *ArXiv:1603.07285 [Cs, Stat]*. <http://arxiv.org/abs/1603.07285>


<https://naokishibuya.medium.com/up-sampling-with-transposed-convolution-9ae4f2df52d0>

Transposed convolutions

- E.g., $i' = 3$, $k = 3$, $s = 1$, $p = 1$



DeconvNet




Semantic Image Segmentation


- Semantic segmentation:
 - Classic techniques using colour
 - Fully Convolutional Network
 - Deconvolution
 - Skip connections: U-Net
- Instance Segmentation
 - Two-stage: Mask-RCNN
 - One-stage: YOLACT

Semantic resolution problem


- Problem: early layers not semantic



Solution 1: Image pyramids




Solution 2: Skip connections




Compute class scores
at multiple layers, then
upsample and add

Solution 2: Skip connections



Red arrows indicate
backpropagation

Skip connections



without skip



with skip


Fully convolutional networks for semantic segmentation. Evan Shelhamer, Jon Long, Trevor Darrell. In *CVPR* 2015

U-Net





Ronneberger, O., Fischer, P., & Brox, T. (2015). U-Net: Convolutional Networks for Biomedical Image Segmentation. *ArXiv:1505.04597 [Cs]*. <http://arxiv.org/abs/1505.04597>

- Network for **semantic segmentation**
- Encoder-decoder-like architecture
- Decoder again contains up-conv layers
- High-resolution feature maps from encoder are combined with feature maps from decoder to improve localization



U-Net



Example result U-Net



Example result U-Net




Semantic Image Segmentation

- Semantic segmentation:
 - Classic techniques using colour
 - Fully Convolutional Network
 - Deconvolution
 - Skip connections: U-Net
- Instance Segmentation
 - Two-stage: Mask-RCNN
 - One-stage: YOLACT

Instance segmentation

Semantic segmentation

Label every pixel, including the background (sky, grass, road)



Do not differentiate between the pixels coming from instances of the same class

Instance segmentation

Label every pixel, including the background (sky, grass, road)



Do not differentiate between the pixels coming from instances of the same class

Do not label pixels coming from uncountable objects (sky, grass, road)



Differentiate between the pixels coming from instances of the same class


Instance segmentation methods

Proposal-based

1. Proposals



2. Assign a class




FCN-based

1. Semantic segmentation



2. Find instances



vs.


Instance segmentation methods

Proposal-based

1. Proposals



2. Assign a class




FCN-based

1. Semantic segmentation



2. Find instances



vs.

Proposal-based methods

Bounding boxes.....

We already know how to obtain those!



Mask R-CNN


Mask R-CNN

He, K., Gkioxari, G., Dollár, P., & Girshick, R. (2017). Mask R-CNN. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2961–2969. <https://doi.org/10.1109/TPAMI.2018.2844175>

- Network for **instance segmentation**
- Extends Faster R-CNN
- Third RoI head that outputs a 14 x 14 **binary mask** for each RoI
- Uses RoIAlign instead of RoIPool
 - Uses bilinear interpolation instead of quantization of box coordinates and when dividing into bins
 - Avoids misalignments introduced by RoIPool


What is Mask-RCNN?

- Starting from the Faster R-CNN architecture




What is Mask-RCNN?

- Faster R-CNN + FCN for segmentation



What is Mask-RCNN?


- Faster R-CNN + FCN for segmentation



Mask loss =
binary cross
entropy per pixel
for the k semantic
classes


He et al. 'Mask R-CNN' ICCV 2017

Mask R-CNN



Detection vs. segmentation

- Detection: for object classification, you require **invariant** representations




Translation invariance: wherever the penguin is in the image, I still want to have “penguin” as my classification output

Detection vs. segmentation

- Detection: for object classification, you require **invariant** representations
- Segmentation: you require **equivariant** representations
 - Translated object → Translated mask
 - Scaled object → scaled mask
 - For semantic segmentation, small objects are less important (less pixels), but for instance segmentation, all objects (no matter the size) are equally important

Mask-RCNN: operations

- What operations are equivariant?




Features extraction = convolutional layers → equivariant

Segmentation head is a fully convolutional network → equivariant

Mask-RCNN: operations

- What operations are equivariant?

Fully connected layers
and global pooling
layers give invariance!




Features extraction = convolutional
layers → equivariant

Segmentation head is a fully
convolutional network → equivariant


Recall: RoI pooling

- Region of Interest Pooling: for every proposal



Recall: RoI pooling


- Let us look at sizes



Mask-RCNN: operations

- Make all operations equivariant


Fully connected layers
and global pooling
layers give invariance!



Exchange RoI pooling by an equivariant operation = RoI Align

RoIAlign

- Erase quantization effects




ROIAlign

To obtain the value
use bilinear
interpolation


Feature map

Each unit is
sampled 4 times

Grid points for
bilinear
interpolation




Mask R-CNN: qualitative results



Mask R-CNN: qualitative results



Mask R-CNN: qualitative results



Mask R-CNN: extended for joints



Model a keypoint's location as a one-hot mask, and adopt Mask R-CNN to predict K masks (which are in the end only 1 pixel), one for each of K keypoint types (e.g., left shoulder, right elbow). This demonstrates the flexibility of Mask R-CNN.


Semantic Image Segmentation

- Semantic segmentation:
 - Classic techniques using colour
 - Fully Convolutional Network
 - Deconvolution
 - Skip connections: U-Net
- Instance Segmentation
 - Two-stage: Mask-RCNN
 - One-stage: YOLACT

Is one-stage vs two-stage also applicable to masks?


One-stage vs two-stage detectors

Faster R-CNN



Slower, but has higher performance


YOLO



Faster, but has lower performance


One-stage vs two-stage instance segmenters

Mask R-CNN



Slower, but has
higher performance

YOLACT



Faster, but has lower
performance

YOLO with masks?


“Boxes are stupid anyway though, I’m probably a true believer in masks except I can’t get YOLO to learn them.”

- Joseph Redmon, YOLOv3


YoLACT*

*You Only Look At CoefficientS

YOLACT: idea




YOLOACT: idea



1) Generate mask prototypes

YOLOACT: idea

2) Generate mask coefficients




1) Generate mask prototypes

YOLOACT: idea


2) Generate mask coefficients

3) Combine (1) and (2)



1) Generate mask prototypes


YOLOACT: backbone



YOLACT: protonet


Generate k prototype masks.

k is not the number of classes, but is a hyperparameter.



YOLACT: protonet

- Fully convolutional network




Similar to the mask branch in Mask R-CNN.


However, no loss function is applied on this stage.

YOLOACT: mask coefficients

Predict a coefficient for every predicted mask.




YOLACT: mask coefficients




The network is similar but shallower than RetinaNet

YOLACT: mask assembly

1. Do a linear combination between the mask coefficients and the mask prototypes.
2. Predict the mask as $M = \sigma(PC^T)$ where P is a ($H \times W \times K$) matrix of prototype masks, C is a ($N \times K$) matrix of mask coefficients surviving NMS, and σ is a nonlinearity.



YOLOACT: loss function




Cross-entropy between the assembled masks and the ground truth, in addition to the standard losses (regression for the bounding box, and classification for the class of the object/mask).

YOLACT: qualitative results




YOLACT: qualitative results



For large objects, the quality of the masks is even better than those of two-stage detectors


So, which segmenter to use?



YOLOACT: improvements


- A specially designed version of NMS, in order to make the procedure faster.
- An auxiliary semantic segmentation loss function performed on the final features of the FPN.
The module is not used during the inference stage.
- D. Boyla et al. “YOLOACT++: Better real-time instance segmentation”. [arXiv:1912.06218](https://arxiv.org/abs/1912.06218) 2019

Example result Instance segmentation (YOLACT++)




171204_pose1 (Min. Cams = 3)


Ground truth




MOBILENETV3+SSD_SMALL



MOBILENETV3+SSD_LARGE



YOLACT++_2D



<https://www.youtube.com/watch?v=rQh2RpEQyKI>

Semantic Image Segmentation

- Semantic segmentation:
 - Classic techniques using colour
 - Fully Convolutional Network
 - Deconvolution
 - Skip connections: U-Net
- Instance Segmentation
 - Two-stage: Mask-RCNN
 - One-stage: YOLACT